

# ‘cs229’—Notes

Malcolm

Started 1st October 2024

# Chapter 1

## Supervised learning

Given a dataset of  $n$  *training examples*  $\{(x^{(i)}, y^{(i)}); i = 1, \dots, n\}$ —a *training set*—where  $\mathbf{x}$  represents the *features* and  $\mathbf{y}$  the “output” or *target* variable we are trying to predict. If not already obvious, we denote the vector space of  $\mathbf{x}$  as  $\mathcal{X}$  and that of the outputs  $\mathbf{y}$  as  $\mathcal{Y}$ .

Our goal is, given a training set, to learn a function  $h : \mathcal{X} \mapsto \mathcal{Y}$  so that  $h(x)$  is a “good” predictor for the corresponding  $y$ . This function  $h$  is called a *hypothesis*.

When trying to predict a continuous target variable, we call this a *regression* problem; whereas when  $y$  can take on only a small number of discrete values we call that a *classification* problem.

### 1.0.1 Linear Regression

Say we decide to approximate  $y$  as a linear function of  $x$ :

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots$$

Where  $\boldsymbol{\theta}$  represents the *parameters/weights* (parametrising the space of linear functions mapping from  $\mathcal{X}$  to  $\mathcal{Y}$ ). We can simplify our notation as such: (by convention letting  $x_0 = 1$ , aptly named the *intercept* term)

$$h(x) = \sum_{i=0}^d \theta_i x_i = \boldsymbol{\theta}^T \mathbf{x}$$

In order to formalise a measure of proximity between the predicted value  $h(x)$  and the target  $y$ , we define a *cost function*:

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

This particular cost function implies an *ordinary least squares* regression model.

### 1.0.2 LMS algorithm

Our cost function  $J(\theta)$  gives us a measure of prediction accuracy. We want to choose  $\theta$  so as to minimise  $J(\theta)$ . Starting with an initial set of  $\theta$ , we need a search algorithm that repeatedly changes  $\theta$  in an attempt to minimise  $J(\theta)$ . Here we consider the *gradient descent* algorithm, which, given some initial  $\theta$ , repeatedly performs the update:

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

Where the update is simultaneously performed for all values of  $j = 0, \dots, d$ .  $\alpha$  is called the *learning rate* (how much we move in the direction the gradient points in).

#### Intuition

Consider attempting to minimise the least mean squares (LMS) cost function for a single training example:

$$\begin{aligned} \frac{\partial}{\partial \theta_j} J(\theta) &= \frac{\partial}{\partial \theta_j} \frac{1}{2} (h_\theta(x) - y)^2 \\ &= 2 \cdot \frac{1}{2} (h_\theta(x) - y) \cdot \frac{\partial}{\partial \theta_j} (h_\theta(x) - y) \\ &= (h_\theta(x) - y) \cdot \frac{\partial}{\partial \theta_j} \left( \sum_{i=0}^d \theta_i x_i - y \right) \\ &= (h_\theta(x) - y) x_j \end{aligned}$$

This gives us the update rule:

$$\theta_j := \theta_j + \alpha \left( y^{(i)} - h_\theta(x^{(i)}) \right) x_j^{(i)}$$

(we use the notation  $a := b$  to denote (in a script) overwriting  $a$  with  $b$ ) Notice the property of the LMS update rule that the magnitude of the update is proportional to the *error* term  $(y^{(i)} - h_\theta(x^{(i)}))$ ; this means that predictions further off the mark result in a greater correction to  $\theta$ .

(next page)

### Batch Gradient Descent

We had the LMS rule for when there was only a single training example. One way to modify this method for a training set of more than one example is the following algorithm:

$$\begin{aligned} &\text{Repeat until convergence } \{ \\ &\quad \theta_j := \theta_j + \alpha \sum_{i=1}^n \left( y^{(i)} - h_{\theta}(x^{(i)}) \right) x_j^{(i)}, \text{ (for every } j) \\ &\} \end{aligned}$$

Written more succinctly ( $\theta_j$  and  $x_j$  as vectors):

$$\theta := \theta + \alpha \sum_{i=1}^n \left( y^{(i)} - h_{\theta}(x^{(i)}) \right) x^{(i)}$$

This method looks at every example in the entire training set on every step, and is called *batch gradient descent*.

### Stochastic gradient descent

Now consider another algorithm:

$$\begin{aligned} &\text{Loop } \{ \\ &\quad \text{for } i = 1 \text{ to } n, \{ \\ &\quad \quad \theta_j := \theta_j + \alpha \left( y^{(i)} - h_{\theta}(x^{(i)}) \right) x_j^{(i)}, \text{ (for every } j) \\ &\quad \} \\ &\} \end{aligned}$$

Written more compactly:

$$\theta := \theta + \alpha \left( y^{(i)} - h_{\theta}(x^{(i)}) \right) x^{(i)}, \text{ (update } \theta \text{ } n \text{ times)}$$

Here we update  $\theta$  for each training example during each run of the training set. This is called *stochastic/incremental gradient descent*.

Whereas batch gradient descent has to scan through the entire training set before taking a single step—which is costly if  $n$  is large—stochastic gradient descent continues to make progress with each example it looks at. Stochastic gradient descent gets  $\theta$  “close” to the minimum much faster than batch gradient descent. Note that “convergence” doesn’t really occur—the parameters  $\theta$  will keep oscillating around the minimum of  $J(\theta)$ . (though most values near minimum would be reasonably good approximations to the true minimum)

### 1.0.3 Gradients of

#### Matrix Derivatives

For a function  $f : \mathbb{R}^{n \times d} \mapsto \mathbb{R}$  mapping from  $n$ -by- $d$  matrices to real numbers, we define the derivative of  $f$  with respect to  $A$  to be

$$\nabla_A f(A) = \begin{bmatrix} \frac{\partial f}{\partial A_{11}} & \cdots & \frac{\partial f}{\partial A_{1d}} \\ \vdots & \ddots & \vdots \\ \frac{\partial f}{\partial A_{n1}} & \cdots & \frac{\partial f}{\partial A_{nd}} \end{bmatrix}$$

#### Gradient of $b^T x$ :

For  $x \in \mathbb{R}^n$  and  $f(x) = b^T x$  for known  $b \in \mathbb{R}^n$ , we have

$$f(x) = \sum_{i=1}^n b_i x_i$$

and so

$$\frac{\partial f(x)}{\partial x_k} = \frac{\partial}{\partial x_k} \sum_{i=1}^n b_i x_i = b_k$$

Consider repeating this for the partial of each element of  $x$ . See that  $\nabla_x b^T x = b$  (analogous to single variable calculus).

#### Gradient of $f(x) = x^T A x$ :

Now consider the quadratic function  $f(x) = x^T A x$  for  $A \in \mathbb{S}^n$  (meaning symmetric). First see that

$$f(x) = \sum_{i=1}^n \sum_{j=1}^n A_{ij} x_i x_j$$

this can be seen from considering that

$$b = Ax \in \mathbb{R}^{n \times 1}$$

can be written as

$$\begin{aligned} b_1 &= \sum_{i=1}^n A_{1i} x_i \\ &\vdots \\ b_n &= \sum_{i=1}^n A_{ni} x_i \end{aligned}$$

(next page)

so we can write

$$\begin{aligned} x^T A x &= x^T b = \sum_{j=1}^n x_j b_j \in \mathbb{R} \\ &= \sum_{j=1}^n x_j \left( \sum_{i=1}^n A_{ji} x_i \right) \end{aligned}$$

Rewriting gives us

$$f(x) = \sum_{i=1}^n \sum_{j=1}^n A_{ij} x_i x_j$$

which was what we wanted. Now we take the partial derivative by considering the terms including  $x_k$  and  $x_k^2$  factors separately:

$$\begin{aligned} \frac{\partial f(x)}{\partial x_k} &= \frac{\partial f(x)}{\partial x_k} \sum_{i=1}^n \sum_{j=1}^n A_{ij} x_i x_j \\ &= \frac{\partial f(x)}{\partial x_k} \left[ \sum_{i \neq k} \sum_{j \neq k} A_{ij} x_i x_j + \sum_{i \neq k} A_{ik} x_i x_k + \sum_{j \neq k} A_{kj} x_k x_j + A_{kk} x_k^2 \right] \\ &= \sum_{i \neq k} A_{ik} x_i + \sum_{j \neq k} A_{kj} x_j + 2A_{kk} x_k \\ &= \sum_{i=1}^n A_{ik} x_i + \sum_{j=1}^n A_{kj} x_j = 2 \sum_{i=1}^n A_{ki} x_i \end{aligned}$$

where the last equality follows since  $A$  is assumed to be *symmetric*. Since

$$\sum_{i=1}^n A_{ki} x_i$$

is just the inner product of a single row of  $A$  and  $x$ —the  $k$ th entry of  $\nabla_x f(x)$  is just the inner product of the  $k$ th row of  $A$  and  $x$ ; therefore

$$\nabla_x x^T A x = 2Ax$$

also analogous to single variable calculus. (next page)

**Hessian**

### 1.0.4 Least Squares matrix representation

Here we consider  $J(\theta)$ —the least squares cost function—in matrix-vectorial notation.

Given a training set, we define the *design matrix*  $\mathbf{X}$  to be the  $n \times d$  matrix ( $n \times d + 1$  if we include the intercept term) that contains the training examples' inputs values in its rows:

$$\mathbf{X} = \begin{bmatrix} -(x^{(1)})^T - \\ -(x^{(2)})^T - \\ \vdots \\ -(x^{(n)})^T - \end{bmatrix}$$

and  $\mathbf{y}$  the  $n$ -dimensional vector containing the target values:

$$\mathbf{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(n)} \end{bmatrix}$$

See that we can represent our least squares function over the entire dataset as